

Dylan Ram

AI Deployment · Agent Systems · Strategy & Enablement

dylanmr96@gmail.com · 916-690-5681 · [linkedin.com/in/dylanram](https://www.linkedin.com/in/dylanram) · github.com/dylan3796/dylan3796.github.io

SUMMARY

Deploys AI into the workflows a go-to-market org runs on — reporting, forecasting, crediting — and sets up how the org adopts it: foundational data, KPI alignment, and self-serve enablement. An operator who builds the systems himself and draws the line where deterministic, tested code must own the outcome. As Databricks' first Partner Strategy & Ops hire, shipped the team's first LLM agents to production and independently built a full agent operating system in Claude Code — skills, hooks, subagents, MCP, and a golden-tested engine that keeps the model out of the money path. Fluent in Python and SQL, and builds with the evals that answer the only question that matters: how do you know it's working.

SKILLS

AI & agents — Claude Code (skills, hooks, subagents, MCP, scoped permissions), agent & LLM workflow design, prompt engineering, evals & golden-test patterns, self-serve enablement, playbooks & standards

Data & engineering — Python (pandas, NumPy), SQL, PySpark, Tableau, Salesforce, medallion architecture, KPI definition, Git/GitHub

EXPERIENCE

Partner Strategy & Ops Manager · **Databricks** *promoted Feb 2023 · first Partner S&O hire* Aug 2021 – Present

- Deployed the partner team's first LLM reporting agents (Newsletter Agent, FAQ Agent) into production — answering partner-metric questions and pushing updates automatically, the team's first agentic reporting motion.
- Set the team's AI strategy and built the foundational data-and-AI layer — a medallion architecture in Spark/SQL with self-serve interfaces so stakeholders answer their own questions.
- Align KPIs across a two-sided marketplace — the partner team, the sales org co-selling through partners, and the partners themselves — the measurement backbone every AI output reports against.
- Designed the partner attribution model (Salesforce/SQL/Spark) and built the team's first revenue forecasting process from scratch — the governed data the agents run on.

SMB Sales Strategy & Operations Analyst · **Salesforce** Jul 2019 – Aug 2021

- Built the territory-carving Python script that encoded the org's guiding principles, personnel, and accounts — the engine behind the annual GTM plan.
- Automated QBR decks, forecast-accuracy tracking, and territory data pulls with Python, APIs, and G Suite.
- Led the business unit's Tableau migration and established data governance for the \$250M AMER SMB Central business.

Business Data Analyst · **CBRE** Sep 2018 – Jul 2019

- Built Python web-scraping tools and managed the product data warehouse; shipped client-facing Tableau dashboards and streamlined the analytics pipeline end to end.

SELECTED PROJECT

Claude Code as an Operating System — a self-built agent platform (public repo) — github.com/dylan3796/anthropic-portfolio

- Architected a five-layer agent operating system: persistent memory, session-boot data hooks, long-horizon plans, invocable skills, and a retro loop that proposes edits to its own setup.
- Drew the AI/deterministic boundary for commissions crediting: an agent turns managers' plain-English coverage into effective-dated rules; a golden-tested Python engine applies them to deals — no LLM ever computes a credited dollar.
- Wrote the eval suite that proves it: golden tests over mid-quarter hires, territory handoffs, split credit, and coverage gaps — an uncredited deal is surfaced, never silently zeroed.
- Built the enablement pattern: the ops team self-served first, then began authoring their own skills — reusable playbooks and standards, the adoption unlock these roles are hired to drive.

EDUCATION

University of California, Santa Barbara — B.A. Economics & Accounting · Dean's Honors