

SKILLS

AI & agents

Claude Code (skills, hooks, subagents, MCP, scoped permissions), agent & LLM workflow design, prompt engineering, evals & golden-test patterns, self-serve enablement, playbooks & standards

Data & engineering

Python (pandas, NumPy), SQL, PySpark, Tableau, Salesforce, medallion architecture, KPI definition, Git/GitHub

LINKS

linkedin.com/in/dylanram
github.com/dylan3796/dylan3796.github.io

EDUCATION

University of California, Santa Barbara

B.A. Economics & Accounting · Dean's Honors

SUMMARY

Deploys AI into the workflows a go-to-market org runs on — reporting, forecasting, crediting — and sets up how the org adopts it: data, KPIs, self-serve enablement. An operator who builds the systems himself and draws the line where deterministic, tested code must own the outcome. As Databricks' first Partner Strategy & Ops hire, shipped the team's first LLM agents to production — reporting agents, partner-recommendation agents for reps' deals, natural-language partner-fit Q&A — and independently built a full agent operating system in Claude Code, with a golden-tested engine that keeps the model out of the money path. Fluent in Python and SQL, and builds with the evals that answer the only question that matters: how do you know it's working.

EXPERIENCE

Partner Strategy & Ops Manager · Databricks

Aug 2021 – Present

promoted Feb 2023 · first Partner S&O hire

- Deployed the partner team's first LLM agents to production: reporting agents (Newsletter, FAQ), partner-recommendation agents for reps' deals, and natural-language partner-fit Q&A.
- Set the team's AI strategy and built the foundational data-and-AI layer — a medallion architecture in Spark/SQL with self-serve interfaces so stakeholders answer their own questions.
- Align KPIs across a two-sided marketplace — the partner team, the sales org co-selling through partners, and the partners themselves — the measurement backbone every AI output reports against.
- Designed the partner attribution model (Salesforce/SQL/Spark) and built the team's first revenue forecasting process from scratch — the governed data the agents run on.

SMB Sales Strategy & Operations Analyst · Salesforce

Jul 2019 – Aug 2021

- Built the territory-carving Python script that encoded the org's guiding principles, personnel, and accounts — the engine behind the annual GTM plan.
- Automated QBR decks, forecast-accuracy tracking, and territory data pulls with Python, APIs, and G Suite.
- Led the business unit's Tableau migration and established data governance for the \$250M AMER SMB Central business.

Business Data Analyst · CBRE

Sep 2018 – Jul 2019

- Built Python web-scraping tools and managed the product data warehouse; shipped client-facing Tableau dashboards and streamlined the analytics pipeline end to end.

SELECTED PROJECT

Claude Code as an Operating System — a self-built agent platform (public repo)

— github.com/dylan3796/anthropic-portfolio

- Architected a five-layer agent operating system: persistent memory, session-boot data hooks, long-horizon plans, invocable skills, and a retro loop that proposes edits to its own setup.
- Drew the AI/deterministic boundary for commissions crediting: an agent turns managers' plain-English coverage into effective-dated rules; a golden-tested Python engine applies them to deals — no LLM ever computes a credited dollar.
- Wrote the eval suite that proves it: golden tests over mid-quarter hires, territory handoffs, split credit, and coverage gaps — an uncredited deal is surfaced, never silently zeroed.
- Built the enablement pattern: the ops team self-served first, then began authoring their own skills — reusable playbooks and standards, the adoption unlock these roles are hired to drive.