

SKILLS

Agents & LLMs

Claude Code (skills, hooks, subagents, MCP, scoped permissions), agent & LLM workflow design, prompt engineering, evals & golden-test patterns

Languages & data

Python (pandas, NumPy), SQL, PySpark, Salesforce APIs, Tableau, Git/GitHub

LINKS

[linkedin.com/in/dylanram](https://www.linkedin.com/in/dylanram)
github.com/dylan3796/anthropic-portfolio

EDUCATION

University of California, Santa Barbara

B.A. Economics & Accounting ·
Dean's Honors

SUMMARY

Operator-engineer who deploys LLM and agent systems into real business workflows — and draws the hard line on where deterministic, tested code must own the outcome. At Databricks, shipped the partner team's first LLM agents to production; independently designed a full agent operating system in Claude Code — skills, hooks, subagents, MCP, and a golden-tested engine that keeps the model out of the money path. Fluent in Python and SQL, and ready to answer the question that matters: how do you know it's working.

EXPERIENCE

Partner Strategy & Ops Manager · Databricks

Aug 2021 – Present

promoted Feb 2023 · first Partner S&O hire

- Deployed the partner team's first LLM reporting agents (Newsletter Agent, FAQ Agent) into production — automated Q&A, insights, and notifications, embedding agentic workflows into how partner metrics reach the org.
- Building the team's foundational data-and-AI layer: a medallion (bronze/silver/gold) architecture in Spark/SQL and self-serve interfaces so stakeholders query answers directly instead of filing requests.
- Aligning KPIs across internal, external, and partner stakeholders as the measurement backbone the AI layer reports against.
- Designed the partner attribution model and built the attribution systems in Salesforce, SQL, and Spark — sourced, influenced, and attributed revenue across a two-sided marketplace; the governed data the agents run on.
- Built the team's first partner-revenue forecasting process where none existed — methodology and cadence designed from scratch.

SMB Sales Strategy & Operations Analyst · Salesforce

Jul 2019 – Aug 2021

- Built the territory-carving Python script that encoded the org's guiding principles, personnel, and accounts — the engine behind the annual GTM plan.
- Automated QBR decks, forecast-accuracy tracking, and territory data pulls with Python, APIs, and G Suite.
- Led the business unit's Tableau migration and established data governance for the \$250M AMER SMB Central business.

Business Data Analyst · CBRE

Sep 2018 – Jul 2019

- Built Python web-scraping tools and managed the product data warehouse; shipped client-facing Tableau dashboards and streamlined the analytics pipeline end to end.

SELECTED PROJECT

Claude Code as an Operating System — a self-built agent platform (public repo, 2026)

— github.com/dylan3796/anthropic-portfolio

- Architected a five-layer agent system: persistent memory, SessionStart hooks that boot each session with live data, long-horizon plans, invocable skills, and a retro loop that proposes edits to its own setup.
- Drew the AI/deterministic boundary for commissions crediting: an agent turns managers' plain-English coverage into effective-dated rules; a golden-tested Python engine applies them to deals — no LLM ever computes a credited dollar.
- Wrote the eval suite that proves it: golden tests over mid-quarter hires, territory handoffs at the boundary, split credit, and coverage gaps — an uncredited deal is surfaced, never silently zeroed.
- Shipped scoped skills, a plan-drafting subagent, least-privilege permissions, and MCP integrations; documented the architecture in a Streamlit walkthrough.